

# Multiple Linear Regression

## Import the required libraries

```
In [77]: 1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5
6 import warnings
7 warnings.filterwarnings('ignore')
```

## Read the data

```
In [78]: 1 housing = pd.read_csv(r"C:\Users\Bhupendra\Desktop\DataCenter\Regressions\Ho
2 housing
```

```
Out[78]:
```

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	hotwaterhe
0	13300000	7420	4	2	3	yes	no	no	
1	12250000	8960	4	4	4	yes	no	no	
2	12250000	9960	3	2	2	yes	no	yes	
3	12215000	7500	4	2	2	yes	no	yes	
4	11410000	7420	4	1	2	yes	yes	yes	
...	...	...	...	...	...	...	...	...	...
540	1820000	3000	2	1	1	yes	no	yes	
541	1767150	2400	3	1	1	no	no	no	
542	1750000	3620	2	1	1	yes	no	no	
543	1750000	2910	3	1	1	no	no	no	
544	1750000	3850	3	1	2	yes	no	no	

545 rows × 13 columns



In [79]: 1 housing.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 545 entries, 0 to 544
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   price                 545 non-null   int64
1   area                 545 non-null   int64
2   bedrooms             545 non-null   int64
3   bathrooms            545 non-null   int64
4   stories              545 non-null   int64
5   mainroad             545 non-null   object
6   guestroom           545 non-null   object
7   basement             545 non-null   object
8   hotwaterheating      545 non-null   object
9   airconditioning      545 non-null   object
10  parking              545 non-null   int64
11  prefarea             545 non-null   object
12  furnishingstatus     545 non-null   object
dtypes: int64(6), object(7)
memory usage: 55.5+ KB
```

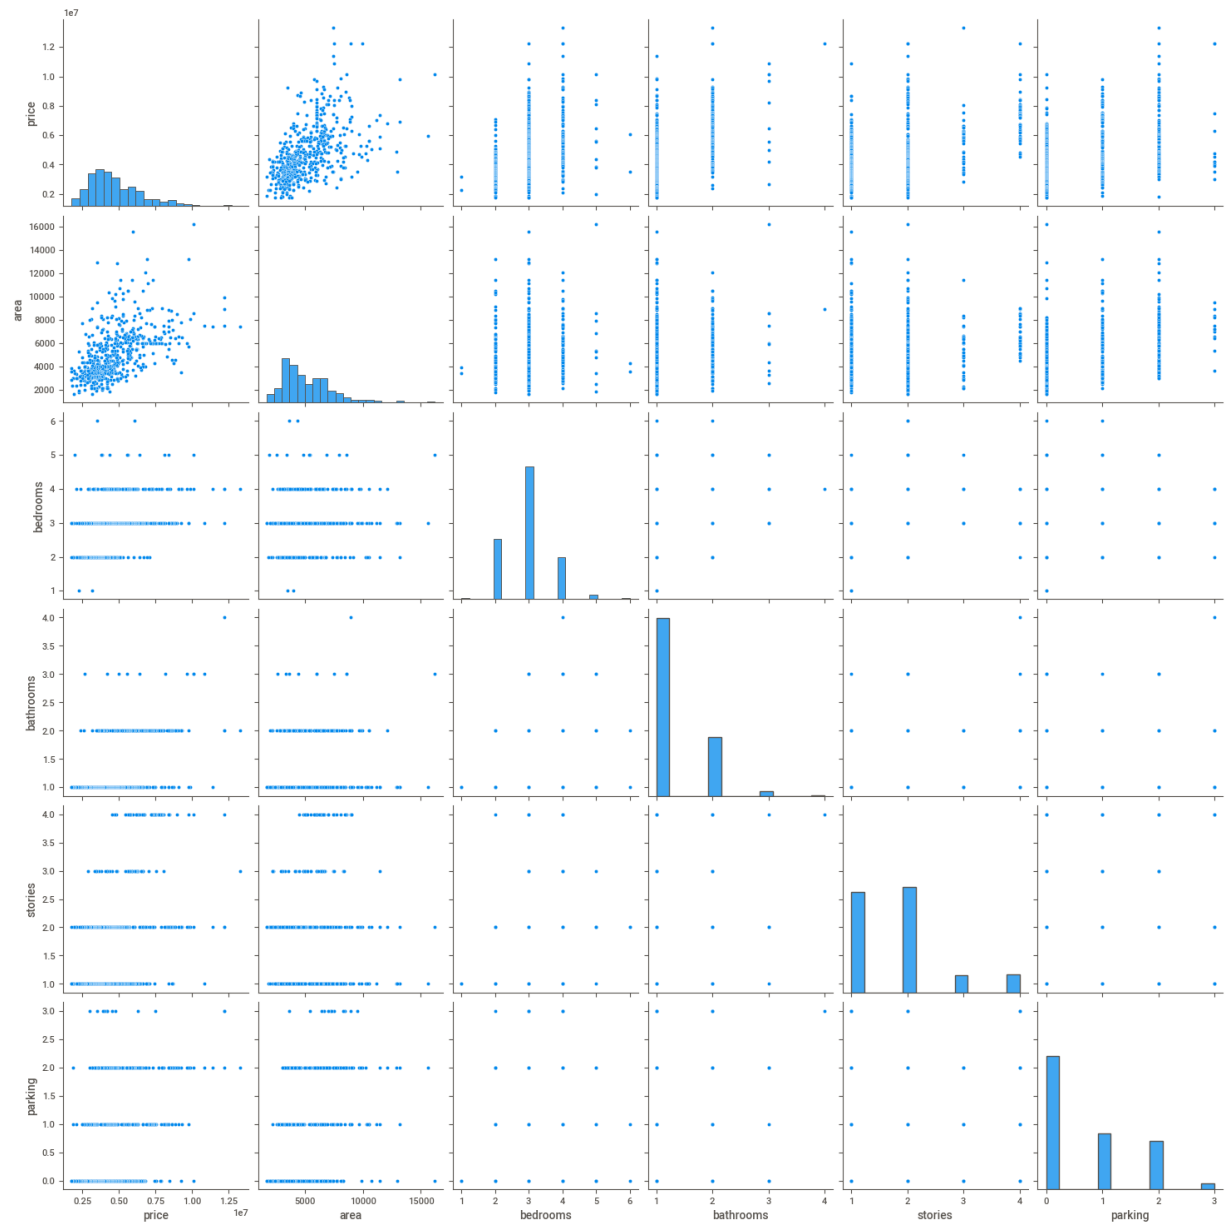
In [80]: 1 housing.describe()

Out[80]:

	price	area	bedrooms	bathrooms	stories	parking
<b>count</b>	5.450000e+02	545.000000	545.000000	545.000000	545.000000	545.000000
<b>mean</b>	4.766729e+06	5150.541284	2.965138	1.286239	1.805505	0.693578
<b>std</b>	1.870440e+06	2170.141023	0.738064	0.502470	0.867492	0.861586
<b>min</b>	1.750000e+06	1650.000000	1.000000	1.000000	1.000000	0.000000
<b>25%</b>	3.430000e+06	3600.000000	2.000000	1.000000	1.000000	0.000000
<b>50%</b>	4.340000e+06	4600.000000	3.000000	1.000000	2.000000	0.000000
<b>75%</b>	5.740000e+06	6360.000000	3.000000	2.000000	2.000000	1.000000
<b>max</b>	1.330000e+07	16200.000000	6.000000	4.000000	4.000000	3.000000

## Visualization

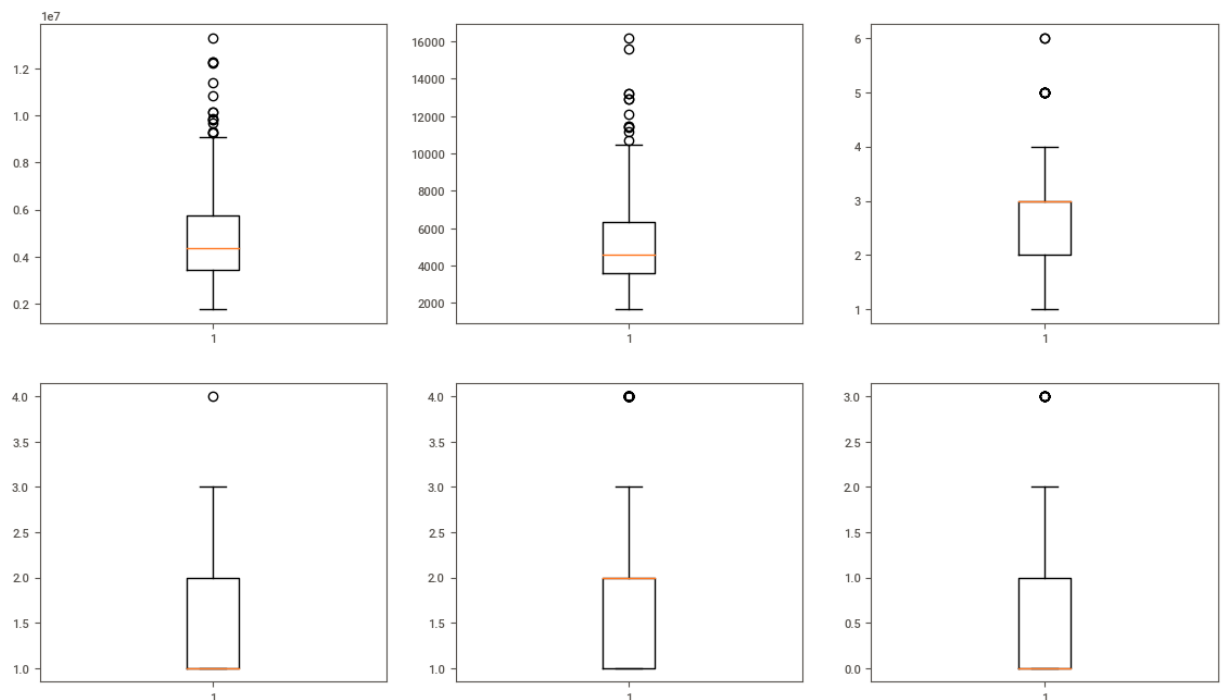
```
In [81]: 1 sns.pairplot(housing)
        2 plt.show()
```



```
In [82]: 1 numeric_cols = housing.select_dtypes(['int', 'float']).columns
2         numeric_cols
```

```
Out[82]: Index(['price', 'area', 'bedrooms', 'bathrooms', 'stories', 'parking'], dtype='object')
```

```
In [83]: 1 fig,ax = plt.subplots(2,3, figsize = (14,8))
2
3 k=0
4 for i in range(0,2):
5     for j in range(0,3):
6         ax[i,j].boxplot(housing[numeric_cols[k]])
7         k+=1
8
9 plt.show()
```



```
In [84]: 1 cat_cols = housing.select_dtypes('object').columns
2         cat_cols
```

```
Out[84]: Index(['mainroad', 'guestroom', 'basement', 'hotwaterheating',
                'airconditioning', 'prefarea', 'furnishingstatus'],
                dtype='object')
```

## SweetViz : automating the data analysis part

```
In [86]: 1 import sweetviz as sv
2 sweet_report = sv.analyze(housing)
3 sweet_report.show_html('housing_analysis.html')
```

| [ 0%] 00:00 ->...

Report housing\_analysis.html was generated! NOTEBOOK/COLAB USERS: the web browser MAY not pop up, regardless, the report IS saved in your notebook/colab file s.

## Label Encoding

```
In [87]: 1 housing.head()
```

```
Out[87]:
```

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	hotwaterheati
0	13300000	7420	4	2	3	yes	no	no	
1	12250000	8960	4	4	4	yes	no	no	
2	12250000	9960	3	2	2	yes	no	yes	
3	12215000	7500	4	2	2	yes	no	yes	
4	11410000	7420	4	1	2	yes	yes	yes	

```
In [88]: 1 housing.mainroad=housing.mainroad.map({'yes':1,'no':0})
2 housing.guestroom=housing.guestroom.map({'yes':1,'no':0})
3 housing.basement=housing.basement.map({'yes':1,'no':0})
4 housing.hotwaterheating=housing.hotwaterheating.map({'yes':1,'no':0})
5 housing.airconditioning=housing.airconditioning.map({'yes':1,'no':0})
6 housing.prefarea=housing.prefarea.map({'yes':1,'no':0})
```

```
In [89]: 1 housing.head()
```

```
Out[89]:
```

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	hotwaterheati
0	13300000	7420	4	2	3	1	0	0	
1	12250000	8960	4	4	4	1	0	0	
2	12250000	9960	3	2	2	1	0	1	
3	12215000	7500	4	2	2	1	0	1	
4	11410000	7420	4	1	2	1	1	1	

```
In [90]: 1 housing.furnishingstatus.value_counts()
```

```
Out[90]: semi-furnished    227
unfurnished    178
furnished    140
Name: furnishingstatus, dtype: int64
```

### Using LabelEncoder class for label encoding

```
In [91]: 1 from sklearn.preprocessing import LabelEncoder
2
3 encoder = LabelEncoder()
4 encoder.fit_transform(housing.furnishingstatus)[:10]
```

```
Out[91]: array([0, 0, 1, 0, 0, 1, 1, 2, 0, 2])
```

```
In [92]: 1 encoder.classes_
```

```
Out[92]: array(['furnished', 'semi-furnished', 'unfurnished'], dtype=object)
```

```
In [93]: 1 housing.furnishingstatus = encoder.fit_transform(housing.furnishingstatus)
```

```
In [94]: 1 housing.head()
```

```
Out[94]:
```

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	hotwaterheati
0	13300000	7420	4	2	3	1	0	0	
1	12250000	8960	4	4	4	1	0	0	
2	12250000	9960	3	2	2	1	0	1	
3	12215000	7500	4	2	2	1	0	1	
4	11410000	7420	4	1	2	1	1	1	

**After converting all categorical columns into numerical ones now our data is ready for modelling**

## Model Training

**train\_test\_split**

```
In [116]: 1 from sklearn.model_selection import train_test_split
2
3 X = housing.drop('price', axis = 1)
4 y = housing.price
5
6 X_train, X_test, y_train, y_test = train_test_split(X,y, test_size = 0.30, r
```

## Model 1

```
In [98]: 1 from sklearn.linear_model import LinearRegression
2
3 # Training
4 model1 = LinearRegression()
5 model1.fit(X_train, y_train)
```

Out[98]: LinearRegression()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

## prediction

```
In [100]: 1 y_pred = model1.predict(X_test)
2 y_pred[:5]
```

Out[100]: array([5407508.87024418, 7097185.46706855, 3055462.44314053,  
4476945.19636315, 3315983.65663579])

```
In [101]: 1 y_test[:5]
```

Out[101]: 316 4060000  
77 6650000  
360 3710000  
90 6440000  
493 2800000  
Name: price, dtype: int64

## Model Evaluation

### r2 score

```
In [102]: 1 model1.score(X_test,y_test)
```

Out[102]: 0.6435419628959107

Model is able to explain **64.35%** variance in the data which is an average score.


## root mean square error

```
In [104]: 1 from sklearn.metrics import mean_squared_error
          2
          3 np.sqrt(mean_squared_error(y_test, y_pred))
```

Out[104]: 1238970.4429194627

## distribution of residuals

```
In [105]: 1 sns.distplot(y_test-y_pred);
```



residuals are normally distributed, satisfying one of the assumptions of OLS(Ordinary Least Square) model

## Improving the model performance

### Model 2



In [107]:

```
1 housing.head()
```

Out[107]:

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	hotwaterheati
0	13300000	7420	4	2	3	1	0	0	
1	12250000	8960	4	4	4	1	0	0	
2	12250000	9960	3	2	2	1	0	1	
3	12215000	7500	4	2	2	1	0	1	
4	11410000	7420	4	1	2	1	1	1	

## Scaling the predictor variables

In [121]:

```
1 from sklearn.preprocessing import StandardScaler, MinMaxScaler
```

In [117]:

```
1 X = housing.drop('price', axis = 1)
2 y = housing.price
```

In [118]:

```
1 X.head(2)
```

Out[118]:

	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	hotwaterheating	aircond
0	7420	4	2	3	1	0	0		0
1	8960	4	4	4	1	0	0		0

In [119]:

```
1 X.columns
```

Out[119]:

```
Index(['area', 'bedrooms', 'bathrooms', 'stories', 'mainroad', 'guestroom',
      'basement', 'hotwaterheating', 'airconditioning', 'parking', 'prefarea',
      'furnishingstatus'],
      dtype='object')
```

In [120]:


```
1 ss = StandardScaler()
2 X = pd.DataFrame(ss.fit_transform(X), columns = X.columns)
3 X.head(2)
```

Out[120]:

	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	hotwaterheating	a
0	1.046726	1.403419	1.421812	1.378217	0.405623	-0.465315	-0.734539	-0.219265	
1	1.757010	1.403419	5.405809	2.532024	0.405623	-0.465315	-0.734539	-0.219265	

## Feature Selection

```
In [128]: 1 plt.figure(figsize = (15,10))
2          sns.heatmap(pd.concat([X,y], axis = 1).corr(), annot = True)
3          plt.show()
```



```
In [140]: 1 X.shape
```

```
Out[140]: (545, 12)
```

```
In [141]: 1 y.shape
```

```
Out[141]: (545,)
```

```
In [139]: 1 from sklearn.feature_selection import SelectKBest
```

```
In [151]: 1 SKB = SelectKBest(k = 8)
2          SKB.fit(X,y)
```

```
Out[151]: SelectKBest(k=8)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [155]: 1 imp_cols = SKB.get_feature_names_out()
          2 imp_cols
```

```
Out[155]: array(['area', 'bedrooms', 'bathrooms', 'stories', 'guestroom',
                'airconditioning', 'prefarea', 'furnishingstatus'], dtype=object)
```

```
In [157]: 1 SKB.n_features_in_
```

```
Out[157]: 12
```

```
In [162]: 1 X.columns
```

```
Out[162]: Index(['area', 'bedrooms', 'bathrooms', 'stories', 'mainroad', 'guestroom',
                'basement', 'hotwaterheating', 'airconditioning', 'parking', 'prefarea',
                'furnishingstatus'],
                dtype='object')
```

```
In [161]: 1 SKB.get_support()
```

```
Out[161]: array([ True,  True,  True,  True, False,  True, False, False,  True,
                False,  True,  True])
```


```
In [163]: 1 imp_cols = X.columns[SKB.get_support()]
          2 imp_cols
```

```
Out[163]: Index(['area', 'bedrooms', 'bathrooms', 'stories', 'guestroom',
                'airconditioning', 'prefarea', 'furnishingstatus'],
                dtype='object')
```

```
In [164]: 1 X = X[imp_cols]
          2 X.head()
```

```
Out[164]:
```

	area	bedrooms	bathrooms	stories	guestroom	airconditioning	prefarea	furnishingstat
0	1.046726	1.403419	1.421812	1.378217	-0.465315	1.472618	1.804941	-1.4062
1	1.757010	1.403419	5.405809	2.532024	-0.465315	1.472618	-0.554035	-1.4062
2	2.218232	0.047278	1.421812	0.224410	-0.465315	-0.679063	1.804941	-0.0916
3	1.083624	1.403419	1.421812	0.224410	-0.465315	1.472618	1.804941	-1.4062
4	1.046726	1.403419	-0.570187	0.224410	2.149083	1.472618	-0.554035	-1.4062



```
In [ ]: 1
```

```
In [ ]: 1
```

```
In [ ]: 1
```

In [ ]:

1

In [ ]:

1

In [ ]:

1

In [122]: 1 X\_train, X\_test, y\_train, y\_test = train\_test\_split(X,y, test\_size = 0.30, r

```
In [123]: 1 from sklearn.linear_model import LinearRegression
2
3 # Training
4 model1 = LinearRegression()
5 model1.fit(X_train, y_train)
```

Out[123]: LinearRegression()

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

In [124]: 1 model1.score(X\_test,y\_test)

Out[124]: 0.6435419628959108

## Model Building

In [167]: 1 X\_train, X\_test, y\_train, y\_test = train\_test\_split(X,y, test\_size = 0.30, r

```
In [168]: 1 model2 = LinearRegression()
2 model2.fit(X_train, y_train)
```

Out[168]: LinearRegression()

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

In [171]: 1 y\_pred = model2.predict(X\_test)

In [172]: 1 model2.score(X\_test,y\_test)

Out[172]: 0.6193144896868267

In [173]: 1 np.sqrt(mean\_squared\_error(y\_test, y\_pred))

Out[173]: 1280383.042743209

In [ ]: 1